



WHAT I BUILT — PRODUCT CLARITY CATALOG



Purpose: one row per active build — what it is, what it does, what purpose it serves, and what it is worth in the 2026 AI market.

Scope: 116 build records with active state (Building or Done) in the live Concepts database, ordered by Concept ID.

Valuation method — read this before using any number

Two numbers are given for every build:

- **As-is** — what the asset is realistically worth today: code/IP only, pre-revenue, no users, private repo, `Build Verified = NO`. This is IP/asset-sale or acqui-hire-contribution value, not company value.
- **Ceiling** — realistic enterprise value *if* the build is runtime-verified, productized, deployed, and reaches early commercial traction, benchmarked against current AI-market comparables in its category.



These are my market estimates, not appraisals and not offers. Nothing here is a verification claim. Zero of these builds currently carry revenue, users, or target-runtime proof, so every As-is figure is a code-asset estimate and every Ceiling figure is conditional.

Category pricing logic used: agent security / agent governance / MCP trust is the highest-multiple category in the current market; AI-provenance and audit infrastructure is second; local/edge inference runtime is third; vertical impact apps and social-good tools price low as assets regardless of engineering quality; research/hardware concepts without physical bench proof price as IP options; and business-model entries are leverage, not sellable products.

Highest-value assets in the estate

Rank	ID	Build	Why it prices highest	As-is	Ceiling
1	T21-100	MCP Server Zero-Trust Sandbox Fabric	MCP is now the default agent-tool protocol and has no trusted isolation layer	\$90K	\$300M
2	T21-094	Prompt-Injection Physics Firewall	Assumes injection succeeds and enforces capability limits outside the model	\$85K	\$300M
3	T21-124	AgentTape Flight Recorder	Deterministic replay is the missing primitive for agent incident response	\$80K	\$250M
4	T21-087	Shadow-AI Connector Inventory	Every enterprise now has unknown MCP/connector surface area	\$75K	\$250M
5	T21-086	Database Agent Safety Kernel	Agents writing to production databases is an unsolved liability	\$75K	\$200M
6	T21-054	Coding-Agent Config Hijack Detector	Real, active attack class against coding agents	\$75K	\$150M
7	T21-083	Malicious Package Family Detector	Cross-ecosystem supply-chain detection is a funded category	\$70K	\$200M
8	T21-140	AI-Native Filesystem Layer	Storage-layer semantics and provenance is platform-level territory	\$70K	\$200M

Rank	ID	Build	Why it prices highest	As-is	Ceiling
9	T21-150	WardenWasm WASI Agent Sandbox	Capability- bounded tool execution with a real standard behind it	\$70K	\$180M
10	T21-131	AgentChaos	Chaos engineering for agents has no established incumbent	\$70K	\$180M

Full catalog — T21-001 to T21-050

ID	Build	What it does	Purpose it serves	As-is	Ceiling
001	Generational Access Vault	Family archive that unlocks messages and recordings on future dates or milestones	Consumer legacy and inheritance of digital memory	\$15K	\$20M
002	Prompt Version Control	Diff, blame, history, and merge for prompts and system instructions	Makes prompt changes auditable engineering artifacts instead of guesswork	\$60K	\$150M
003	AI Cost-of- Ownership Auditor	Profiles real workloads and compares API spend against local/open- model operation	Gives CFO- grade evidence for the cloud- to-local migration decision	\$40K	\$60M
004	Endangered Language Fine- Tune Kit	Community pipeline to adapt a small model from elder recordings	Language preservation owned by the community, not a vendor	\$10K	\$5M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
005	Hallucination Ledger	Versioned public registry of model failures by version, date, and prompt pattern	CVE-style historical record for AI errors	\$20K	\$40M
006	Elder Digital Guardian	On-device scam-call and scam-text detection with family alerting	Protects older users from fraud without cloud surveillance	\$30K	\$75M
007	Dataset Lineage Chain	Merkle evidence chain proving which datasets entered each fine-tune step	Training-data accountability and regulatory defensibility	\$50K	\$80M
008	Home Compute Mesh Scheduler	Pools a user's idle devices into a fault-tolerant local inference cluster	Turns owned hardware into usable AI capacity	\$45K	\$100M
010	AI Estate Executor	Executes predefined shutdown, transfer, or memorialization of accounts and agents after death	Digital estate settlement for an agent-owning world	\$15K	\$25M
013	Encrypted Personal Knowledge Vault	Fully local semantic index over personal files and messages	Private memory and retrieval that never leaves the device	\$30K	\$60M
014	Idle-Device Inference Marketplace	Rents spare compute for inference jobs with execution verification	Monetizes idle hardware with proof the job actually ran	\$20K	\$80M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
015	Verified Model Provenance Marketplace	Model listings carrying verifiable lineage, data hashes, and benchmark evidence	Buying models without trusting the seller's claims	\$35K	\$70M
019	Auto-Generated Model Nutrition Labels	Scans a training pipeline and emits standardized transparency disclosures	Automates AI-transparency compliance paperwork	\$20K	\$25M
020	Reproducibility Pinning Service	Freezes weights, quantization, runtime, and seeds so outputs can be reproduced	Makes AI outputs auditable after the fact	\$55K	\$90M
021	Degraded-Network Model Router	Falls back from cloud to local model on connectivity loss, then resyncs	AI that keeps working on bad networks	\$45K	\$60M
022	Cross-Vendor Agent Interop Bus	Shared schema and transport so agents from different vendors coordinate	Prevents one model vendor from owning the coordination layer	\$60K	\$200M
027	Dialect-Faithful Transcription	Transcribes preserving regional and minority dialect features	Accuracy and dignity for speakers standard ASR erases	\$35K	\$50M
028	Camera-Firmware Provenance Signing	Signs media at capture before software can alter it	Origin proof for media independent of any platform	\$60K	\$150M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
029	Remittance Verification	Local pattern checks that flag suspicious recipient changes on family transfers	Fraud protection for informal money movement without a bank cloud	\$20K	\$30M
032	Solar AI-in-a-Box Kits	Documented solar Pi-class local-LLM bundles for low-connectivity sites	Deployable AI where there is no grid or network	\$15K	\$20M
033	Water Sensor and Local AI Alert Mesh	Solar sensor network doing local inference for contamination and failure alerts	Water safety monitoring with no cloud subscription	\$25K	\$40M
034	Small-Church/Nonprofit AI Ops Kit	Local AI toolkit for scheduling, bulletins, follow-up, and admin	Capable operations for organizations with no IT staff	\$10K	\$10M
037	Personal Baseline Drift Tracker	Tracks changes in a person's own typing and speech patterns over time	Early self-observed change detection without diagnosing	\$15K	\$30M
038	Model Weight Tamper-Evidence	Verifies downloaded open weights match the published artifact	Stops silently poisoned or swapped model weights	\$50K	\$80M
039	Solo-Creator Chargeback Shield	Fraud-pattern detection and chargeback defense sized for individual sellers	Protects solo revenue from fraud losses enterprises absorb easily	\$25K	\$40M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
040	AI Proof-of-Presence	Attests with location, sensors, and time that a device action happened as claimed	Trustworthy evidence for autonomous device activity	\$65K	\$150M
041	Offline Crop Diagnosis	Phone-camera local model identifying crop disease or deficiency offline	Agronomy advice with no connectivity requirement	\$30K	\$50M
042	Offline Two-Way Sign Language Translator	Local speech-to-sign and sign-to-speech translation	Communication access where no interpreter exists	\$40K	\$80M
043	Reentry Communication Assistant	Helps drafting letters, reading paperwork, and tracking deadlines on restricted devices	Reduces reentry failure caused by paperwork and comms barriers	\$15K	\$25M
044	Gig-Worker Earnings Copilot	Reads pay statements and contracts and flags underpayment or mismatch	Wage recovery leverage on the worker's side	\$25K	\$50M
045	Federated Fine-Tune Co-op	Shared fine-tuning where raw data never leaves devices and value is shared	Collective model improvement without data surrender	\$45K	\$120M
046	On-Device Biometric Auth	Drop-in local face or voice authentication	Biometric login without sending biometrics to a cloud vendor	\$55K	\$100M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
047	Energy-Metered Local Inference Logger	Measures real energy per inference run and reports cost and carbon	Actual AI energy accounting instead of estimates	\$35K	\$60M
049	Agent Handoff Receipt	Verifiable receipt when one agent passes work to another	Chain of custody and accountability between agents	\$55K	\$120M
050	Poison Detector for Solo Fine-Tuners	Scans training data for backdoors and poisoning patterns	Training-data safety for developers without a security team	\$45K	\$90M

Full catalog — T21-051 to T21-100

ID	Build	What it does	Purpose it serves	As-is	Ceiling
051	Frugal Code Healer	Coding agent that tests its own output and learns cheaper reliable tool-use patterns	Reliable coding assistance on modest hardware	\$40K	\$80M
052	Polyglot Tool-Call Contract Verifier	Enforces formal cross-language contracts on agent tool calls at invocation	Stops agents from calling tools outside their allowed effects	\$70K	\$200M
053	Power-Thrifty Self-Improving Translator	On-device translation that improves low-resource pairs within device limits	Translation quality gains without cloud or new hardware	\$30K	\$50M
054	Coding-Agent Config Hijack Detector	Detects symlink, environment, and config hijacks before tool execution	Blocks a live attack class against AI coding agents	\$75K	\$150M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
055	EdgeHeal Recursive Farmer	Solar edge AI diagnosing crop disease and improving from verified outcomes	Field-learning agronomy with rollback and drift control	\$30K	\$50M
056	Agentic Termination Guardian	Synthesizes and verifies termination invariants and halts unbounded agent loops	Prevents runaway agent loops mathematically, not heuristically	\$65K	\$120M
057	On-Device Research Intern	Local research agent for hypothesis generation and experiment design	Research throughput without cloud lab tooling	\$20K	\$30M
058	Failure-Mining Infrastructure Agent	Monitors pumps and solar gear and improves rules from real technician repairs	Predictive maintenance that learns from actual field fixes	\$40K	\$80M
059	Recursive Skill Weaver for Artisans	Observes task sequences, tests improved procedures, shares over Bluetooth mesh	Skill transfer and improvement in trades with no documentation	\$15K	\$20M
062	Village Voice Loop	SMS and USSD AI for feature phones in local languages with community correction	AI access for users who have no smartphone or data	\$20K	\$40M
063	Recursive Curriculum Composer	Generates local-context lessons and evolves pedagogy on detected plateaus	Curriculum that adapts to the actual learners	\$20K	\$40M
065	Live Exploitability	Simulates exploit paths through	Tells you if agent-	\$70K	\$150M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
	Path Simulator	dependency graphs an agent modified	introduced code is actually exploitable		
066	Outcome-Driven Micro-Insurance AI	Lightweight underwriting and claims that learn from real payouts and detect bias	Insurable coverage for populations traditional underwriting excludes	\$20K	\$60M
067	Community-Verified Fact Weaver	Builds local knowledge bases and isolates contested low-confidence claims	Local truth maintenance without central editorial control	\$20K	\$30M
070	Sensor-Fusion Self-Healer	Improves multimodal fusion and reroutes inference when a sensor dies	Edge perception that degrades gracefully instead of failing	\$35K	\$70M
071	Adaptive Toolsmith Agent	Designs and iterates physical tools and jigs from real performance feedback	Local manufacturing improvement loop	\$20K	\$30M
072	Energy-Aware Recursive Controller	Micro-grid controller improving forecasting under tight power budgets	Reliable power control with failing sensors and actuators	\$40K	\$90M
073	Soil and Seed Recursive Advisor	Learns planting and fertilizer strategy from farmer experiments under climate drift	Agronomic decisions that survive changing conditions	\$25K	\$50M
074	Agent Sandbox Escape Proof Generator	Minimizes an escape into a	Turns a vague agent incident into a	\$60K	\$100M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
		reproducible proof	patchable root cause		
075	Behavioral Baseline for Excessive Agency	Learns normal agent action distributions and blocks anomalous high-agency sequences	Detects an agent doing far more than its task requires	\$70K	\$180M
076	Local Market Price Healer	SMS-first commodity-price forecasting under sparse data and shocks	Pricing power for sellers with no market data	\$15K	\$25M
077	Recursive Craft Skill Library	Extracts reusable skills from video and sensors and repairs incomplete encodings	Preserves and transmits craft technique	\$15K	\$20M
078	Unbounded Consumption Rate-Limit Synthesizer	Generates and enforces adaptive token and compute budgets for agent workflows	Stops runaway agent spend	\$55K	\$100M
079	HITL Validation Accelerator for AI PRs	Builds minimal high-signal review checklists and differential tests from AI changes	Makes human review of AI code fast and targeted	\$50K	\$120M
080	Agent Containment Escape Replay Engine	Deterministically replays trajectories that crossed a sandbox boundary	Red-team training and patch generation from real escapes	\$55K	\$90M
081	Adaptive Micro-Credit Scorer	Auditable edge credit model refining itself from repayment outcomes	Lending decisions for informal borrowers with drift detection	\$25K	\$60M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
082	Mesh-Healing Logistics Agent	Decentralized last-mile routing exchanging partial models over local radio	Delivery coordination with no reliable network	\$25K	\$50M
083	Malicious Package Family Detector	Links related malicious packages across PyPI, npm, and Docker after repackaging	Catches supply-chain attack families, not single samples	\$70K	\$200M
084	Semantic Snowflaking Lineage Reconstructor	Reconstructs functional ancestry across AI-generated forks despite divergent history	Provenance and license enforcement when Git history lies	\$45K	\$80M
085	Self-Patching Micro-Clinic	Portable offline triage box improving thresholds from verified outcomes	Frontline clinical triage with no connectivity	\$25K	\$40M
086	Database Agent Safety Kernel	Forces database agents through verified side-effect-bounded interfaces	Prevents agents from destroying production data	\$75K	\$200M
087	Shadow-AI Connector Inventory and Risk Scorer	Discovers and scores every MCP server and connector reachable by the org	Visibility into unsanctioned AI attack surface	\$75K	\$250M
088	Swarm-Healing Environmental Monitor	Cheap sensor swarm that detects failed peers and compensates collaboratively	Environmental data that stays reliable as nodes die	\$30K	\$60M
089	Human Authorship	Cryptographically separates human	Evidence for copyright,	\$50K	\$100M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
	Attestation Protocol	conception from AI assistance	licensing, and inventorship		
090	EOL Exposure Shadow Runtime	Probes end-of-life components for new exploitability and proposes isolation	Manages risk in software that can never be upgraded	\$50K	\$90M
091	Drift-Resistant Water Sentinel	Ultra-low-power water sensing that compensates for failed nodes synthetically	Irrigation and water anomaly detection with sparse intervention	\$25K	\$45M
092	Community Knowledge Compactor	Captures oral knowledge into a local graph and quarantines low-consensus entries	Institutional memory for communities with no written record	\$25K	\$40M
093	Review-Queue Load Balancer with Slop Pre-Filter	Routes AI contributions by predicted review cost and filters low-signal ones	Protects maintainer time from AI submission floods	\$40K	\$80M
094	Prompt-Injection Physics Firewall	External control plane enforcing hard capability limits assuming injection succeeds	The only defense posture that survives prompt injection	\$85K	\$300M
095	Intent-Preserving Spec-to-Agent Compiler	Compiles requirements into constrained agent plans that hold intent	Stops agency drift between spec and execution	\$60K	\$150M
096	Code Novelty and Hallucination Differential	Measures novelty and factual grounding of AI code against the live OSS corpus	Tells you whether AI code is original, copied, or invented	\$45K	\$80M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
097	Self-Repairing Language Bridge	Tiny multilingual local model improving from community correction	Usable AI for rare and Indigenous languages	\$25K	\$40M
098	Sparse-Data Disease Sentinel	Outbreak detection from sparse reports with reporting-bias compensation	Public-health early warning where surveillance is thin	\$30K	\$60M
099	Auto-Evolving School Tutor	SMS, WhatsApp, and offline tutor that replaces ineffective teaching paths	Personalized instruction on the devices students already have	\$30K	\$80M
100	MCP Server Zero-Trust Sandbox Fabric	Treats every MCP tool call as untrusted and attenuates capability automatically	The trust layer the MCP ecosystem currently lacks	\$90K	\$300M

Full catalog — T21-102 to T21-182

ID	Build	What it does	Purpose it serves	As-is	Ceiling
102	Self-Explaining Runtime	Translates services, decisions, jobs, and failures into plain-English live state	Operability without reading source or dashboards	\$40K	\$80M
103	Distributed Local Memory Fabric	Private memory layer shared across a user's own devices	Continuous AI memory that never centralizes	\$50K	\$120M
116	Portable Identity Across Networks	Identity that survives network changes without re-login or setup	Removes per-network identity friction	\$45K	\$100M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
124	AgentTape Deterministic Flight Recorder	Records and deterministically replays agent actions, tool calls, and state	Incident response, debugging, and liability evidence for agents	\$80K	\$250M
126	TensorWire Zero-Copy Local AI IPC	Moves tensors between local processes without copying	Removes the main local multi-process inference bottleneck	\$55K	\$100M
128	Customer-Owned Compute	Runs the product on compute the customer controls	Sells software into buyers who refuse vendor cloud	\$20K	Model, not product
131	AgentChaos	Injects controlled failures to test agent recovery and containment	Proves agent systems survive real failure	\$70K	\$180M
136	Global Surplus Bandwidth Market Layer	Pools and exchanges unused bandwidth and compute	Monetizes idle network capacity	\$15K	\$60M
137	Sunlight-Driven Diffractional Metasurface Processors	Passive optical classification from ambient light with no digital vision pipeline	Zero-power sensing and classification	\$20K	IP option
139	Cross-Device Application Runtime	Moves active tasks across devices without restarting the workflow	Continuity of work across hardware	\$50K	\$120M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
140	AI-Native Filesystem Layer	Adds semantics, provenance, relationships, and local indexing at the storage layer	Makes the filesystem itself AI-usable	\$70K	\$200M
143	Adaptive UI SDK	Interface runtime that adapts to device, task, and context	One UI codebase across very different surfaces	\$30K	\$60M
144	Zero-Cloud Phone Intelligence Stack	On-device speech, vision, and automation with no cloud dependence	A private phone AI stack	\$60K	\$200M
145	Hardware Autopilot	Detects hardware and tunes model and runtime choices automatically	Removes manual performance tuning per device	\$45K	\$90M
146	Interactive Investor Demos	Demos that expose working behavior and evidence instead of slides	Sales and fundraising proof surface	\$10K	Internal leverage
148	One-Binary Infrastructure Appliance	Single Rust binary launching database, eventing, inference, workers, and dashboard	Full stack with no orchestration expertise	\$55K	\$150M
150	WardenWasm WASI Agent Sandbox	Runs agent tools as isolated WASI components with explicit capabilities	Bounded tool execution on an open standard	\$70K	\$180M
151	Local Voice Computer	Offline voice control with local reasoning, actions, and	Voice computing that is not a cloud microphone	\$35K	\$80M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
		persistent memory			
152	SCX-Infer AI-Aware Linux Scheduler	sched_ext scheduler accounting for inference workload behavior	Real throughput gains on shared inference hosts	\$60K	\$120M
153	Recursive Build Repair Engine	Consumes build failures, patches, rebuilds, and retests in a bounded loop	Ends the manual build-fix cycle	\$50K	\$100M
154	AgentKernel	Minimal deterministic runtime over task, capability, process, evidence, and result	The primitive layer agent platforms are missing	\$60K	\$150M
156	Local AI Operating Runtime	Single local control plane for models, hardware, and devices	One place to operate local AI	\$55K	\$150M
160	White-Label Product Factory	Deploys multiple branded variants from one shared codebase	Multiplies revenue per unit of engineering	\$15K	Model, not product
161	Self-Porting Compute Identity	Portable digital self and compute context that follows the user	Identity and context independent of any platform	\$40K	\$100M
162	ForgeCell Universal WASM Plugin Runtime	OCI-distributed WASI runtime exposing plugins over HTTP, CLI, and MCP	One safe plugin system across interfaces	\$55K	\$120M
165	Domain Leasing and Licensing	Turns owned domains and	Recurring income from	\$5K	Depends on portfolio

ID	Build	What it does	Purpose it serves	As-is	Ceiling
		product surfaces into leasable assets	dormant domain inventory		
166	Acquisition Bundles	Packages software, domains, proof, and operations into acquisition-ready units	Makes the estate legible to buyers	\$5K	Packaging, not product
167	ColdStart Edge Model Loader	Optimized model loading for constrained edge and fast local startup	Removes cold-start latency on edge inference	\$45K	\$80M
168	Autonomous Civic Mesh Internet	Self-healing city-scale mesh network operating autonomously	Connectivity independent of carriers	\$30K	\$150M capex-bound
169	Micro-SaaS Domain Farm	Portfolio of narrow utilities from shared components at near-zero marginal cost	Many small cash-flow products from one base	\$10K	Cash-flow play
170	MeshInfer Edge Inference Mesh	Coordinates inference across heterogeneous edge hardware with no vendor lock	Uses whatever accelerators exist	\$55K	\$120M
171	ToolFuzz MCP Protocol Fuzzer	Fuzzes MCP tool interfaces to expose malformed-input, state, and security faults	Hardens the MCP ecosystem before attackers do	\$50K	\$100M
172	Sovereign Agent Runtime	Runs bounded, recoverable, evidenced agent processes under local control	Agents you own and can prove behavior for	\$60K	\$150M

ID	Build	What it does	Purpose it serves	As-is	Ceiling
173	Universal Edge Compute Fabric	Pools phones, laptops, VMs, and SBCs into one compute fabric	Capacity from hardware you already own	\$55K	\$150M
175	Salvaged Analog Compute Tile	Reconfigurable analog solver from reclaimed op-amps and microcontrollers	Ultra-low-power optimization and inference from e-waste	\$15K	IP option
182	Vibration-Powered Event Tag	Harvester powers a tag only when a specified impact or vibration occurs	Battery-free event detection on machines and cargo	\$15K	IP option

Portfolio-level read



Sum of As-is line items: roughly **\$4.6M**. **Realistic estate value as a single block sale today: \$1.5M–\$3.5M** — buyers discount heavily for unverified runtime, overlapping scope, and zero users. **Concentrated ceiling:** about **\$2.5B–\$3B** of the total ceiling sits in roughly 15 agent-security and agent-governance assets.

What the numbers say about strategy:

- The estate's value is **not** distributed evenly across 116 builds. Around 15 builds carry most of it: T21-100, 094, 124, 087, 086, 054, 083, 140, 150, 131, 075, 052, 065, 022, 144.
- The social-impact and vertical builds are strong work with low asset prices. They matter for mission and narrative, not for valuation.
- Research and hardware entries such as T21-137, 175, and 182 price as IP options until physical bench proof exists.
- Business-model entries such as T21-160, 165, 166, and 169 are leverage on the rest of the estate and should not be counted as products.

The single biggest value unlock is not another build. Moving **Build Verified** from NO to YES on the top 15, with target-runtime proof and one public runnable demo each, is worth more than the remaining 100 builds combined at current pricing.

Reconstructed catalog — T21-021 through T21-034

T21-021 — Degraded-Network Model Router

- **Plain English:** Automatically falls back from cloud inference to a smaller local model when connectivity degrades, then synchronizes state after recovery.
- **How it works:** Checks connection quality and model availability, then routes work to the best reachable local or remote model while preserving a degraded/offline path.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** Field teams, rural/remote deployments, edge-AI operators, disaster-response systems, and apps that cannot assume stable internet.
- **How to run:** bash [boot021.sh](#)
- **State:** Building · Build Verified NO

T21-022 — Cross-Vendor Agent Interop Bus

- **Plain English:** Shared message schema and transport allowing agents from different model vendors to coordinate without one vendor owning the pipe.
- **How it works:** Normalizes messages and tool calls across incompatible agent systems so they can exchange work through one typed bus.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** Enterprises and developers running agents from multiple vendors or frameworks.
- **How to run:** bash [boot022.sh](#)
- **State:** Building · Build Verified NO

T21-027 — Dialect-Faithful Transcription

- **Plain English:** Transcription system that preserves regional and minority dialect features instead of normalizing them to standard spelling.
- **How it works:** Transcribes speech while preserving dialect-specific wording and pronunciation choices instead of automatically converting them to standardized language.
- **What it is:** Standalone Product · Commercial path: Enterprise License
- **Who needs it:** Media teams, researchers, community archives, courts/interview workflows, and speakers whose dialects are flattened by normal transcription.
- **How to run:** bash [boot027.sh](#)
- **State:** Building · Build Verified NO

T21-028 — Camera-Firmware Provenance Signing

- **Plain English:** Signs media at capture time before downstream software modifies it, providing origin evidence independent of the hosting platform.
- **How it works:** Signs firmware and capture provenance so an image or video can be tied back to a known device/software state and checked for tampering.
- **What it is:** Feature / Component · Commercial path: Data / Registry
- **Who needs it:** Camera/device makers, journalists, evidence workflows, industrial inspection teams, and organizations validating captured media.
- **How to run:** bash [boot028.sh](#)
- **State:** Building · Build Verified NO

T21-029 — Remittance Verification for Diaspora Transfers

- **Plain English:** Local pattern-checking tool for informal family money transfers that can flag suspicious recipient changes without depending on a bank cloud.
- **How it works:** Tracks the promised transfer, fees, exchange terms, and recipient confirmation to produce a verifiable record of what was sent and received.
- **What it is:** Standalone Product · Commercial path: Enterprise License
- **Who needs it:** People sending money across borders, remittance services, community organizations, and recipients who need proof that the intended amount arrived.
- **How to run:** bash [boot029.sh](#)
- **State:** Building · Build Verified NO

T21-032 — Solar AI-in-a-Box Kits

- **Plain English:** Documented solar-powered Pi-class local-LLM bundles for translation, literacy, triage, and other low-connectivity deployments.
- **How it works:** Packages local AI compute with solar-power constraints so the system can schedule workloads around available energy and continue without cloud dependence.
- **What it is:** Standalone Company · Commercial path: Appliance / Kit
- **Who needs it:** Remote schools, farms, clinics, field teams, disaster-response groups, and communities without reliable grid power.
- **How to run:** bash [boot032.sh](#)
- **State:** Building · Build Verified NO

T21-033 — Water Sensor + Local AI Alert Mesh

- **Plain English:** Solar sensor network using local inference to detect contamination or equipment failure without a cloud subscription.
- **How it works:** Combines distributed water sensors with local analysis and mesh alerts so abnormal readings can trigger action even when internet service is unavailable.
- **What it is:** Standalone Company · Commercial path: Appliance / Kit
- **Who needs it:** Farms, rural communities, facilities, utilities, environmental teams, and disaster-response groups.
- **How to run:** bash [boot033.sh](#)
- **State:** Building · Build Verified NO

T21-034 — Small-Church/Nonprofit AI Ops Kit

- **Plain English:** Simple local AI operations toolkit for scheduling, bulletins, follow-up, and administration where there is no dedicated IT staff.
- **How it works:** Bundles common administrative AI workflows into a controlled local toolkit for communications, records, scheduling, and repetitive office work.
- **What it is:** Standalone Company · Commercial path: Appliance / Kit
- **Who needs it:** Small churches, nonprofits, community organizations, and volunteer-run operations with little technical staff.
- **How to run:** bash [boot034.sh](#)
- **State:** Building · Build Verified NO

Reconstructed catalog — T21-037 through T21-044

T21-037 — Personal Baseline Drift Tracker

- **Plain English:** Opt-in local tool that tracks changes in a person's own typing or speech patterns over time without making a diagnosis.
- **How it works:** Builds a personal baseline from repeated local observations and flags meaningful drift from that person's history instead of comparing everyone to one population average.
- **What it is:** Standalone Product · Commercial path: Enterprise License
- **Who needs it:** Individuals, caregivers, wellness/research teams, and longitudinal-monitoring products that care about change from a person's own normal pattern.
- **How to run:** bash [boot037.sh](#)

- **State:** Building · Build Verified NO

T21-038 — Model Weight Tamper-Evidence

- **Plain English:** Cryptographic verification that downloaded open-model weights match the originally published artifact.
- **How it works:** Hashes and records model-weight artifacts so later copies can be checked against the expected identity and modification history.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** Model distributors, enterprises, labs, edge deployments, and security teams that need to know whether a model artifact changed.
- **How to run:** bash [boot038.sh](#)
- **State:** Building · Build Verified NO

T21-039 — Solo-Creator Chargeback Shield

- **Plain English:** Fraud-pattern detection and chargeback defense sized for individual online sellers rather than enterprise merchants.
- **How it works:** Preserves delivery, approval, access, and transaction evidence so a solo seller can answer disputed-payment claims with a coherent record.
- **What it is:** Standalone Product · Commercial path: Enterprise License
- **Who needs it:** Independent creators, freelancers, small digital sellers, and payment-support teams.
- **How to run:** bash [boot039.sh](#)
- **State:** Building · Build Verified NO

T21-040 — AI Proof-of-Presence

- **Plain English:** Cryptographic attestation using location, sensors, and timestamps to prove an autonomous device action occurred where and when claimed.
- **How it works:** Produces signed or hashed runtime evidence tying an AI execution to a model, device/runtime context, and time instead of relying on a screenshot or claim.
- **What it is:** Standalone Company · Commercial path: Data / Registry
- **Who needs it:** Organizations and users that need evidence that a particular AI model/runtime was actually present and executing on a specific device or environment.
- **How to run:** bash [boot040.sh](#)
- **State:** Building · Build Verified NO

T21-041 — Offline Crop Diagnosis

- **Plain English:** Phone-camera local model for identifying crop disease or nutrient deficiency without a data connection.
- **How it works:** Analyzes crop observations locally, classifies supported crop problems, and returns agricultural decision-support without requiring a cloud connection.
- **What it is:** Standalone Company · Commercial path: Appliance / Kit
- **Who needs it:** Farmers, extension workers, agronomy teams, remote agricultural programs, and field users without dependable connectivity.
- **How to run:** bash [boot041.sh](#)
- **State:** Done · Build Verified NO

T21-042 — Offline Two-Way Sign Language Translator

- **Plain English:** Local speech-to-sign and sign-to-speech translation for places where a human interpreter is unavailable.
- **How it works:** Runs the two-way translation workflow locally so signing and speaking users can communicate when an interpreter or network connection is unavailable.
- **What it is:** Standalone Company · Commercial path: Appliance / Kit
- **Who needs it:** Deaf and hard-of-hearing users, hearing users communicating with signers, schools, clinics, public-service desks, field teams, and accessibility developers.
- **How to run:** bash [boot042.sh](#)
- **State:** Done · Build Verified NO

T21-043 — Reentry Communication Assistant

- **Plain English:** Helps incarcerated people and families draft letters, understand paperwork, and track deadlines within restricted-device environments.
- **How it works:** Organizes and assists with reentry-related communication, forms, obligations, and service interactions so important steps are easier to understand and complete.
- **What it is:** Standalone Company · Commercial path: Appliance / Kit
- **Who needs it:** People returning from incarceration, reentry organizations, case workers, employers/service providers, and families navigating required communication.
- **How to run:** bash [boot043.sh](#)
- **State:** Done · Build Verified NO

T21-044 — Gig-Worker Earnings Copilot

- **Plain English:** Reads gig pay statements and contracts and flags possible underpayment or contract mismatches from the worker's side.
- **How it works:** Ingests contract/pay information, calculates what should have been paid, compares it with actual pay, and produces mismatch evidence from the worker's side.
- **What it is:** Standalone Company · Commercial path: SaaS
- **Who needs it:** Delivery, rideshare and freelance workers, worker advocates, unions, legal-aid teams, payroll reviewers, and contractor-compensation auditors.
- **How to run:** bash [boot044.sh](#)
- **State:** Done · Build Verified NO

Reconstructed catalog — T21-045 through T21-053

T21-045 — Federated Fine-Tune Co-op

- **Plain English:** Participants contribute compute and private data to shared fine-tuning without raw data leaving devices and share in resulting usage value.
- **How it works:** Participants train locally on private data, submit compatible model updates rather than raw records, aggregate those updates, and track contribution value.
- **What it is:** Service / Business Model · Commercial path: Service Contract
- **Who needs it:** Private AI co-ops, research groups, community model projects, enterprises with siloed data, and regulated teams that cannot pool raw datasets.
- **How to run:** bash [boot045.sh](#)
- **State:** Done · Build Verified NO

T21-046 — On-Device Biometric Auth

- **Plain English:** Drop-in local face or voice authentication for applications that cannot or do not want to use cloud biometric services.
- **How it works:** Enrolls biometric templates locally, validates new samples against server-controlled thresholds, binds templates to subjects, and enforces lockout/failure rules.
- **What it is:** Standalone Product · Commercial path: Appliance / Kit
- **Who needs it:** Apps, local services, device owners, enterprise edge deployments, and privacy-sensitive products needing a biometric factor without a cloud identity provider.
- **How to run:** bash [boot046.sh](#)
- **State:** Done · Build Verified NO

T21-047 — Energy-Metered Local Inference Logger

- **Plain English:** Measures real energy consumption per local inference run and produces reporting data for cost and carbon accounting.
- **How it works:** Measures inference activity together with energy use and records an auditable per-job cost/energy history.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** Local-AI operators, edge-device fleets, sustainability teams, FinOps teams, and researchers measuring actual inference cost.
- **How to run:** bash [boot047.sh](#)
- **State:** Done · Build Verified NO

T21-049 — Agent Handoff Receipt

- **Plain English:** Verifiable audit receipt when one agent passes work to another, recording what was transferred, what changed, and who accepted it.
- **How it works:** Creates a durable receipt when work, context, authority, or artifacts move between agents so the handoff can be audited later.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** Teams running multiple agents, autonomous workflows, regulated automation, and anyone who needs proof of what one agent passed to another.
- **How to run:** bash [boot049.sh](#)
- **State:** Done · Build Verified NO

T21-050 — Poison Detector for Solo Fine-Tuners

- **Plain English:** Training-data scanner for injected backdoors and poisoning patterns sized for individual developers and small teams.
- **How it works:** Scans fine-tuning inputs and training signals for suspicious patterns so poisoned examples can be isolated before they contaminate a model.
- **What it is:** Standalone Company · Commercial path: SaaS
- **Who needs it:** Independent model builders, small AI teams, dataset curators, and security teams fine-tuning from untrusted or mixed-source data.
- **How to run:** bash [boot050.sh](#)
- **State:** Building · Build Verified NO

T21-051 — Frugal Code Healer

- **Plain English:** Lightweight coding agent that tests its own generated code, rewrites prompts/scaffolding/verification strategies, and learns cheaper, more reliable tool-use patterns on modest hardware.
- **How it works:** Inspects failing code/tests locally, proposes bounded repairs, reruns verification, and keeps the repair loop resource-conscious.
- **What it is:** Standalone Product · Commercial path: Enterprise License
- **Who needs it:** Developers, maintainers, small teams, and local coding-agent users who want repair without expensive API loops.
- **How to run:** bash [boot051.sh](#)
- **State:** Building · Build Verified NO

T21-052 — Polyglot Agent Tool-Call Contract Verifier

- **Plain English:** Formal cross-language contracts for agent tool calls verified at invocation time to constrain arguments, effects, and audit requirements.
- **How it works:** Checks tool schemas, arguments, responses, and behavior against the same contract across different language implementations.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** Agent-platform developers, MCP/tool authors, enterprise integration teams, and anyone maintaining tools across several programming languages.
- **How to run:** bash [boot052.sh](#)
- **State:** Building · Build Verified NO

T21-053 — Power-Thrifty Self-Improving Translator

- **Plain English:** On-device translation system that improves low-resource language pairs from user interactions while adapting latency, energy use, pruning, and architecture to device limits.
- **How it works:** Performs translation locally while tracking resource use and using feedback to improve within an energy budget rather than constantly retraining in the cloud.
- **What it is:** Standalone Product · Commercial path: Appliance / Kit
- **Who needs it:** Mobile/edge translation users, field teams, language projects, and device makers operating under tight battery/compute limits.
- **How to run:** bash [boot053.sh](#)
- **State:** Building · Build Verified NO

Reconstructed catalog — T21-054 through T21-065

T21-054 — AI-Coding-Agent Config Hijack Detector (GitSpawn-Class)

- **Plain English:** Static and dynamic analysis for symbolic-link, environment, and configuration-based coding-agent hijacks before tool execution.
- **How it works:** Inspects agent-relevant configuration and repository changes for instructions or settings that could silently redirect a coding agent's behavior.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** Developers, security teams, CI/CD owners, and organizations allowing coding agents to read repository configuration.
- **How to run:** bash [boot054.sh](#)
- **State:** Done · Build Verified NO

T21-055 — EdgeHeal Recursive Farmer

- **Plain English:** Offline edge AI on cheap solar-powered hardware that diagnoses crop disease and recursively improves a TinyML model from verified local outcomes and farmer feedback, with rollback and drift handling.
- **How it works:** Collects local farm signals, evaluates problems at the edge, recommends/executes bounded responses, and feeds verified outcomes back into the next decision cycle.
- **What it is:** Standalone Company · Commercial path: SaaS
- **Who needs it:** Farms, greenhouse operators, agricultural co-ops, extension programs, and remote agritech deployments.
- **How to run:** bash [boot055.sh](#)
- **State:** Building · Build Verified NO

T21-056 — Agentic Termination Guardian with Formal Loop Invariants

- **Plain English:** Runtime that synthesizes and verifies mathematical termination invariants for agent loops across languages and stops loops that violate bounded execution.
- **How it works:** Defines invariants around agent loops and stops execution when termination/safety conditions are violated rather than allowing unbounded self-calling.
- **What it is:** Feature / Component · Commercial path: Enterprise License

- **Who needs it:** Agent-platform builders, autonomous-workflow operators, safety teams, and systems where runaway loops create cost or operational risk.
- **How to run:** bash [boot056.sh](#)
- **State:** Building · Build Verified NO

T21-058 — Failure-Mining Infrastructure Agent

- **Plain English:** Monitors pumps, solar installations, or small machines, diagnoses failures, and improves preventive rules from successful technician repairs.
- **How it works:** Ingests failures and logs, groups recurring causes, extracts repair knowledge, and turns repeated infrastructure breakage into reusable fixes.
- **What it is:** Standalone Company · Commercial path: SaaS
- **Who needs it:** DevOps, SRE, platform engineering, and teams with recurring build/deployment incidents.
- **How to run:** bash [boot058.sh](#)
- **State:** Done · Build Verified NO

T21-059 — Recursive Skill Weaver for Artisans

- **Plain English:** Mobile/edge system that observes artisan task sequences, invents and tests improved procedural skills, rolls back failed techniques, and shares local knowledge over Bluetooth mesh.
- **How it works:** Captures practical skill steps, links them into reusable learning paths, and refines the skill graph from verified practice and feedback.
- **What it is:** Standalone Company · Commercial path: SaaS
- **Who needs it:** Tradespeople, craftspeople, vocational educators, apprenticeships, and organizations preserving hands-on expertise.
- **How to run:** bash [boot059.sh](#)
- **State:** Done · Build Verified NO

T21-062 — Village Voice Loop

- **Plain English:** SMS/USSD AI for feature phones that answers in local languages, improves from community corrections, and isolates hallucinated or culturally mismatched responses.
- **How it works:** Collects community observations/requests, groups and routes them, records responses, and closes the loop so people can see what happened after they spoke.
- **What it is:** Standalone Company · Commercial path: Appliance / Kit

- **Who needs it:** Local communities, neighborhood organizations, civic groups, nonprofits, and service providers needing a structured local feedback loop.
- **How to run:** bash [boot062.sh](#)
- **State:** Building · Build Verified NO

T21-063 — Recursive Curriculum Composer

- **Plain English:** Generates educational micro-content for local contexts and languages, detects learning plateaus or cultural mismatch, and iteratively evolves pedagogy.
- **How it works:** Builds a curriculum from goals and learner evidence, then adjusts later lessons according to what was actually mastered or missed.
- **What it is:** Standalone Company · Commercial path: SaaS
- **Who needs it:** Teachers, tutors, training teams, homeschoolers, workforce programs, and adaptive-learning products.
- **How to run:** bash [boot063.sh](#)
- **State:** Building · Build Verified NO

T21-065 — Live Exploitability Path Simulator for Agent-Touched Code

- **Plain English:** Simulates realistic exploit paths through dependency graphs modified or recommended by an AI agent.
- **How it works:** Maps code changed by an agent into reachable attack paths and simulates exploitability so reviews focus on concrete reachable risk rather than generic warnings.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** Security teams, developers using coding agents, AppSec teams, and CI pipelines that need to know whether an AI change creates a real attack path.
- **How to run:** bash [boot065.sh](#)
- **State:** Done · Build Verified NO

Reconstructed catalog — T21-070 through T21-077

T21-070 — Sensor-Fusion Self-Healer

- **Plain English:** Low-power multimodal edge system that improves camera/sensor fusion and reroutes inference through remaining modalities when a sensor fails.

- **How it works:** Cross-checks several sensor streams, identifies a drifting or failed source, degrades safely, and reweights/recoveries the system instead of trusting one bad reading.
- **What it is:** Standalone Company · Commercial path: Appliance / Kit
- **Who needs it:** Robotics, IoT, edge-device fleets, industrial systems, autonomous equipment, and teams operating multiple imperfect sensors.
- **How to run:** bash [boot070.sh](#)
- **State:** Done · Build Verified NO

T21-071 — Adaptive Toolsmith Agent

- **Plain English:** Edge AI that designs and iterates simple physical tools or jigs, learns from real-world performance feedback, and evolves new search strategies after failures.
- **How it works:** Detects repeated capability gaps, builds or adapts bounded tools for those gaps, tests them, and makes successful tools available to later runs.
- **What it is:** Standalone Company · Commercial path: SaaS
- **Who needs it:** Agent-platform developers, automation teams, and power users whose agents repeatedly encounter missing tool capabilities.
- **How to run:** bash [boot071.sh](#)
- **State:** Building · Build Verified NO

T21-072 — Energy-Aware Recursive Controller

- **Plain English:** Micro-grid or solar-home controller that improves load forecasting and control under tight power budgets and works around sensor or actuator failures.
- **How it works:** Chooses control/compute actions using both task goals and available energy, then adjusts future decisions from measured power use.
- **What it is:** Standalone Company · Commercial path: Appliance / Kit
- **Who needs it:** Battery/solar edge systems, robotics, mobile devices, remote installations, and compute fleets with strict energy budgets.
- **How to run:** bash [boot072.sh](#)
- **State:** Building · Build Verified NO

T21-073 — Soil & Seed Recursive Advisor

- **Plain English:** Edge AI that learns local planting and fertilizer strategies from farmer experiments and adapts to climate drift using uncertainty estimates and synthetic scenario testing.

- **How it works:** Combines soil/seed observations with local recommendations and feeds measured crop outcomes back into later planting/advice decisions.
- **What it is:** Standalone Company · Commercial path: White-Label
- **Who needs it:** Farmers, agronomists, seed programs, community agriculture projects, and remote field teams.
- **How to run:** bash [boot073.sh](#)
- **State:** Building · Build Verified NO

T21-074 — Agent Sandbox Escape Proof Generator

- **Plain English:** Automatically minimizes an agent sandbox escape into a reproducible proof to accelerate root-cause analysis and patching.
- **How it works:** Runs adversarial containment tests and emits evidence showing which sandbox boundaries held or failed instead of merely claiming isolation.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** Agent-runtime builders, security teams, sandbox vendors, and enterprises running untrusted autonomous code.
- **How to run:** bash [boot074.sh](#)
- **State:** Building · Build Verified NO

T21-075 — Behavioral Baseline Generator for Excessive Agency

- **Plain English:** Learns normal agent action distributions by task class and flags or blocks statistically anomalous high-agency sequences.
- **How it works:** Learns the normal action/tool-call envelope of an agent and flags deviations that indicate the agent is attempting more autonomy than its baseline permits.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** AI safety teams, agent operators, enterprises, and developers monitoring autonomous systems for unexpected expansion of authority.
- **How to run:** bash [boot075.sh](#)
- **State:** Building · Build Verified NO

T21-076 — Local Market Price Healer

- **Plain English:** SMS-first local commodity-price forecasting system that learns from actual market outcomes and adapts to sparse data and sudden shocks with uncertainty-aware strategies.

- **How it works:** Collects local price observations, detects stale/outlier values, reconciles trusted updates, and maintains a healthier shared local price picture.
- **What it is:** Standalone Company · Commercial path: SaaS
- **Who needs it:** Small merchants, farmers, local marketplaces, co-ops, and communities where pricing data is fragmented or stale.
- **How to run:** bash [boot076.sh](#)
- **State:** Building · Build Verified NO

T21-077 — Recursive Craft Skill Library

- **Plain English:** Captures artisan techniques from video and sensors, extracts reusable skills, and repairs incomplete skill encodings by requesting targeted clarification or generating variants.
- **How it works:** Stores craft procedures as linked skills and improves the library from verified attempts, corrections, and newly learned subskills.
- **What it is:** Standalone Company · Commercial path: Appliance / Kit
- **Who needs it:** Craft workers, makers, vocational schools, apprenticeships, repair communities, and organizations preserving practical know-how.
- **How to run:** bash [boot077.sh](#)
- **State:** Building · Build Verified NO

Reconstructed catalog — T21-078 through T21-085

T21-078 — Unbounded Consumption Rate-Limit Synthesizer

- **Plain English:** Generates and enforces adaptive token and compute budgets for agent workflows based on observed behavior.
- **How it works:** Observes how an agent or service consumes tools/resources and generates enforceable rate limits that bound runaway usage without relying on a hard-coded one-size limit.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** Agent-platform teams, API owners, cloud/AI cost-control teams, and safety engineers.
- **How to run:** bash [boot078.sh](#)
- **State:** Building · Build Verified NO

T21-079 — Human-in-the-Loop Validation Accelerator for AI PRs

- **Plain English:** Creates minimal high-signal review checklists and differential tests from AI-generated changes plus the original intent.
- **How it works:** Prepackages the evidence a human reviewer needs—changed behavior, tests, risk areas, and validation checkpoints—so AI-authored PRs can be reviewed faster without removing human approval.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** Software teams using AI-generated pull requests, code reviewers, CI teams, and engineering managers.
- **How to run:** bash [boot079.sh](#)
- **State:** Building · Build Verified NO

T21-080 — Agent Containment Escape Replay Engine

- **Plain English:** Records and deterministically replays trajectories that cross sandbox boundaries for red-team training and patch generation.
- **How it works:** Records a containment escape attempt as a replayable trace so the same sequence can be rerun against repaired sandboxes and future versions.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** AI security teams, sandbox developers, agent-platform operators, and incident-response teams.
- **How to run:** bash [boot080.sh](#)
- **State:** Building · Build Verified NO

T21-081 — Adaptive Micro-Credit Scorer

- **Plain English:** Auditable edge credit model for informal lending that refines itself from repayment outcomes and detects concept drift or demographic shifts.
- **How it works:** Updates small-loan risk decisions from local repayment/outcome history while preserving an auditable scoring trail instead of using a static generic score.
- **What it is:** Standalone Company · Commercial path: SaaS
- **Who needs it:** Community lenders, microfinance programs, co-ops, small-business lending tools, and financial inclusion projects.
- **How to run:** bash [boot081.sh](#)
- **State:** Building · Build Verified NO

T21-082 — Mesh-Healing Logistics Agent

- **Plain English:** Decentralized last-mile routing and demand system whose nodes exchange partial models over Bluetooth/Wi-Fi Direct and isolate compromised or stale models.
- **How it works:** Replans routes/tasks when nodes, links, vehicles, or depots fail and propagates the revised plan across the remaining mesh.
- **What it is:** Standalone Company · Commercial path: SaaS
- **Who needs it:** Logistics teams, field operations, disaster-response networks, warehouses, and distributed delivery systems.
- **How to run:** bash [boot082.sh](#)
- **State:** Building · Build Verified NO

T21-083 — Cross-Ecosystem Malicious Package Family Detector

- **Plain English:** Identifies related malicious packages across ecosystems such as PyPI, npm, and Docker even after repackaging, using structural and behavioral fingerprints.
- **How it works:** Correlates package behavior and metadata across ecosystems to identify related malicious package families instead of scanning each registry in isolation.
- **What it is:** Feature / Component · Commercial path: Enterprise License
- **Who needs it:** Software supply-chain security teams, package registries, enterprise AppSec, and developers consuming npm/PyPI/Cargo/etc.
- **How to run:** bash [boot083.sh](#)
- **State:** Building · Build Verified NO

T21-084 — Semantic Snowflaking Lineage Reconstructor

- **Plain English:** Reconstructs functional ancestry across AI-generated forks even when source text and Git history diverge, using semantic and behavioral fingerprints.
- **How it works:** Reconstructs likely ancestry among semantically similar artifacts so teams can trace fragmented derivatives back toward their source lineage.
- **What it is:** Feature / Component · Commercial path: Data / Registry
- **Who needs it:** Data/ML governance teams, analysts, investigators, and organizations whose datasets/models have been copied and transformed repeatedly.
- **How to run:** bash [boot084.sh](#)
- **State:** Building · Build Verified NO

T21-085 — Self-Patching Micro-Clinic

- **Plain English:** Portable offline diagnostic box using a repurposed phone and sensors; improves triage thresholds from verified outcomes and monitors false-positive and false-negative drift.
- **How it works:** Monitors the local clinic software stack for known operational failures, applies bounded verified repairs, and preserves rollback/audit evidence rather than silently patching.
- **What it is:** Standalone Company · Commercial path: Appliance / Kit
- **Who needs it:** Small clinics, remote-health deployments, field-health organizations, and operators maintaining local clinical software with limited IT support.
- **How to run:** bash [boot085.sh](#)
- **State:** Building · Build Verified NO